AI Growth Strategy for Riverbend Physical Therapy

The Vcrsoft Way: Understand → Diagnose → Identify AI solutions → Prioritize → Deliver. Business Snapshot is step 1, Growth Diagnosis is step 2, Opportunities are step 3, Prioritized Top 3 is step 4, and the Executive Summary, 30-Day Plan, Risks & Metrics are step 5.

Physical Therapy Clinic · Riverbend Physical Therapy · 2026-11-01 · Report #47

Executive Summary

Riverbend Physical Therapy is a two-site outpatient practice (Asheville and Hendersonville, NC) collecting about $75,000 per month across roughly 650 completed visits, with a stated goal of $1.25-1.3 million in collected revenue within twenty-four months. The interview evidence points to a specific growth constraint: demand is healthy, but usable capacity is not — prime-time hours run at 92-95 percent while midday sits at 55-60 percent, 11-12 percent of booked visits are lost to no-shows and late cancels, and the ~35-patient waitlist is mostly prime-time-only so midday slots rarely refill. Three opportunities emerged from deep probing, and the owner tested and reordered the draft ranking himself: (1) waitlist and cancellation texting tied to the WebPT calendar, (2) a measured restart of the clinic's own AI scribe pilot with a named note auditor, and (3) referral source cleanup with a monthly referrer report and a physician feedback loop. All three fit the owner's stated $500-800 monthly tool budget; the two costlier risks are unverified WebPT integrations and the office manager's 4-5 hour weekly implementation bandwidth, so the plan sequences the rollouts one at a time. The Hendersonville expansion decision (adjacent-suite lease, deadline end of Q1 2027) is gated on the owner's own criteria — a signed clinician hire and a $85,000 monthly collection run rate — not on the tools in this report; the tools protect and grow the revenue that must clear those gates.

Business Snapshot

Riverbend Physical Therapy operates two clinics: the flagship in Asheville (opened 2014, five treatment rooms plus an open gym area) and a second site in Hendersonville (opened 2021, three rooms). The clinical team is nine clinicians — seven physical therapists and two PTAs, a composition the owner corrected during draft review after the first draft miscounted it — led by owner Dana Whitmore, who still treats about 25 hours per week of a ~55-hour week. Support staff are office manager Jessica (Asheville), one front-desk person per site, and a remote

verification specialist (Marta). Services are orthopedic and manual therapy (~70 percent of visits), sports medicine, pelvic health (one therapist, Nadia, with a 4-6 week wait), vestibular rehab (owner-only), and dry needling. Payer mix is roughly 45 percent commercial, 30 percent Medicare, 15 percent workers comp/auto, and 10 percent cash pay (pelvic wellness packages, dry needling, and a $99/month 'Fit-for-Life' membership with 11 members). Volume is ~650 completed visits per month (~430 Asheville, ~220 Hendersonville) and ~90 new evaluations per month, at an average net of about $115 per visit (~$105 Hendersonville vs ~$120 Asheville — not yet reconciled against payer and contract mix). Collected revenue is ~$75,000 per month (~$900,000 annualized) at a 13-15 percent margin.

Growth Aspiration

The owner's stated goal is to grow collected revenue from ~$900,000 to $1.25-1.3 million within twenty-four months (~850 visits per month), built on two moves: expanding the Hendersonville suite into the adjacent 1,200-square-foot space (two more rooms, buildout estimated $85,000-110,000) before the landlord's offer expires at the end of Q1 2027, and building pelvic health into a two-therapist line. A third location north of Asheville is an aspiration for later, not a plan. The owner has set explicit gates for the expansion: $85,000 monthly collections for 90 days, a signed clinician hire before buildout starts, prime-time utilization above 85 percent (already met), and four months of operating cash in reserve (~$120,000; currently ~$95,000, so this gate is unmet). The owner also wants his own treating hours down from ~25 to ~12 per week, with freed time going to referrer relationships, hiring, and expansion preparation.

Growth Diagnosis

Growth at Riverbend is gated by usable capacity and owner time, not by demand. The evidence: prime-time slots run 92-95 percent full while midday runs 55-60 percent; ~35 patients wait, about 80 percent of them for prime hours only, so midday cancellations almost never refill (front desk phone-tag fills maybe half of prime-time cancels same-day, midday rarely); 11-12 percent of booked visits are lost to no-shows and late cancels against email-only reminders; and pelvic health turns away 4-6 new patients a month against a 4-6 week wait — unserved demand, not missing demand. Every filled cancellation, in the owner's words, is 'cash in the door this week' at no incremental clinical cost. Underneath the capacity problem sit two staff-time drains that make the constraint expensive: documentation consumes 45-60 minutes per clinician per day (the team's number-one morale complaint), and pre-visit verification takes 20-30 minutes per new patient in batched payer portal sessions, producing ~3 wrong-benefit quotes last quarter and 2-3 late workers-comp authorizations a month that surface as rework

and denials (eligibility/authorization is a top denial category alongside late-signed plans of care and unit coding at the outsourced biller, 6-8 percent first-pass denials, ~34 days AR). Finally, the owner treating ~25 hours a week caps the growth work his own expansion gates require — hiring, referrer relationships, and buildout prep happen only after 6 p.m. Alternatives were tested against this diagnosis: weak overall demand is rejected by the transcript (referral inflow steady, waitlists present), though Hendersonville's referral base is genuinely thinner (5-6 active referrers vs ~12-15 in Asheville) — a site-specific demand risk the referral-cleanup opportunity addresses; front-end conversion is largely healthy (~90 percent of completed evaluations start a plan of care), with the leakage concentrated in after-hours callback speed and insurance confusion rather than in the evaluation itself; and a pricing/mix play (premium cash programs) remains the owner's own hypothesis, unproven and repeatedly deferred for lack of time. The disciplined provisional reading: fix capacity leakage and staff-time drains first, and treat Hendersonville demand development as the parallel track that protects the expansion thesis.

Sales & Growth Engine

New patients arrive roughly half through physician referrals (~45-50 percent; 12-15 actively sending practices, led by Asheville Orthopedic Associates, Blue Ridge Spine Center, and two primary care groups), ~20-25 percent from past patients and word of mouth, ~18 percent from Google search and the website, ~8 percent from direct-access self-referrals (mostly pelvic and athletes; North Carolina allows direct access but Medicare and most commercial plans still tie care to a referring provider), and the remainder from attorneys and employers on workers comp/auto. Two structural weaknesses sit inside this engine. First, the WebPT referral source field is free text, so referrer performance cannot be reported reliably — a sports-medicine referrer who dropped from 6-8 patients a month to 1-2 went unnoticed for about three months until a hallway conversation at a conference. Second, physician communication is one-directional: discharge summaries are faxed ad hoc and Medicare plans of care go back automatically, but there are no systematic mid-episode progress updates. Reviews are strong where they exist (148 reviews at 4.8 in Asheville; 41 at 4.7 in Hendersonville) but only about a third of discharged patients are asked, and reactivation is unowned — one 2024 email to ~300 past patients produced 4 bookings and was never repeated; ~1,800 past-patient emails sit unused in WebPT. Marketing spend is ~$1,400 per month of Google Ads (owner says the figure is approximate), Hendersonville-only since January 2025 after the Asheville campaign was cut for lack of attribution, plus ~$350 in sponsorships; the ads may drive 10-15 calls a month but cannot be tied to booked evaluations. Downstream, ~60-65 percent of inquiries become scheduled evaluations and ~90 percent of completed evaluations start a plan of care (~9 visits ortho, ~8 pelvic, ~6 vestibular); about two-thirds of plans complete, with drop-off clustering at

visits 4-6 when patients feel better, copays accumulate, and rides get hard — and drift is spotted reactively (no-shows, front-desk flags) because nobody runs WebPT's completion report on a schedule.

Current AI Use

Two uses were disclosed in the interview. Owner Dana Whitmore personally uses ChatGPT for marketing copy and job postings — nothing involving patient data. The clinic ran a six-week AI scribe pilot (two clinicians, March-April 2026, $99 per clinician per month, a BAA signed) that saved 35-45 minutes per clinician per day but required 5-8 minutes of editing per note and a copy-paste transfer into WebPT; it had no note-audit step and no baseline measurements, paused when the owner stepped away for a family illness in May 2026, and was never restarted. Administratively the practice runs on WebPT (records, scheduling, billing workflow) with Google Workspace, basic VoIP phones with voicemail, a spreadsheet-and-phone-calls waitlist, and WebPT automated emails as the only reminders. No other AI tools were disclosed in this interview.

Opportunities

Automated waitlist and cancellation filling with two-way texting

Problem: Prime-time slots (6:30-9 a.m., 4-7 p.m.) run 92-95 percent full and midday runs 55-60 percent; a ~35-person waitlist is ~80 percent prime-time-only; cancellations are worked by front-desk phone calls against a spreadsheet, filling maybe half of prime-time cancels same-day and almost no midday slots; and 11-12 percent of booked visits are lost to no-shows and late cancels with email-only reminders.

Why it matters: The owner named fill rate as the fastest lever on the P&L ('a filled cancellation is cash in the door this week with zero clinical cost'). Recovering lost prime slots and refilling midday cancellations attacks revenue directly at ~$115 average net per visit, supports the $85,000/month collection gate for the expansion, and requires no clinician behavior change.

AI approach: A two-way texting layer integrated with the WebPT calendar: when a cancellation opens a slot, matched waitlist patients (by time-of-day preference) get an automatic text offer; confirmations and reminder sequences run by text with reply-based confirm/cancel handling; rule setup and monitoring stay with the office manager.

How: Cancellation appears in WebPT; the tool texts the next matched waitlist patients; first reply wins the slot and the calendar updates (or front desk confirms if integration is manual); reminders go out at 48 and 24 hours and the morning of the visit; unfilled slots escalate to a phone call. Requires the consent script at intake before the first text.

Impact: Mechanism, not a promise: the clinic books ~650 visits against 11-12 percent leakage and a half-filled same-day prime recovery — each recovered prime slot is ~$115 of same-week revenue, and cutting no-shows 3 points is roughly 20 visits a month across both sites. The 90-day keep-it standard the owner set: at least 60 percent of open slots filled within two business days and no-shows down 3 points.

Effort: Within the owner's $500-800/month budget. Hard dependencies to verify before contracting: whether the chosen tool integrates with the WebPT calendar (the owner made this a condition — 'it has to actually talk to WebPT's calendar'); consent/opt-in wording at intake; Jessica's setup time kept inside her 4-5 weekly hours. Kill criteria the owner set: hot opt-out rates, or calendar sync that fights WebPT.

Validation: Baseline fill rate and no-show rate for two weeks before launch; then track weekly: percent of open slots filled within 2 business days (target >=60 percent), no-show +

> late-cancel rate (target -3 points from the 11-12 percent range), and opt-out rate. Jessica pulls WebPT reports; owner reviews monthly.

Measured AI scribe restart with a named note auditor

Problem: Documentation takes 45-60 minutes per clinician per day after the last patient and is the team's number-one morale complaint. The clinic's own six-week, two-clinician pilot (spring 2026, $99/clinician/month) already proved 35-45 minutes of daily savings, but it ran without an owner, without a note-audit step, and without any baseline measurement, used manual copy-paste into WebPT, and lapsed in May 2026 when the owner stepped away — the same no-owner pattern as the membership program.

Why it matters: Nine clinicians recovering even 30 minutes a day frees roughly a clinician-day per week of capacity, and the owner links his own freed hours to the exact work his expansion gates require (referrer outreach, hiring, buildout prep). It also addresses the stated morale complaint, which matters for retaining clinicians in a 3-5-month hiring market.

AI approach: Ambient AI scribe redrafting visit notes, with the treating clinician reviewing and signing every note — restart as a bounded, measured program rather than the unmeasured first pilot.

How: First answer the integration question: the owner never asked WebPT whether a scribe integration or API path exists — the first pilot was copy-paste from the scribe portal. Then restart with two clinicians (owner and Nadia), Mark — the skeptic — converted into the quality owner auditing a weekly sample of five notes, a documented baseline week for documentation minutes, and the owner's keep-it metrics checked at 30 days. Sequenced after the texting rollout so Jessica's rollout bandwidth is never double-booked.

Impact: The pilot's own numbers — 35-45 minutes saved per clinician per day, 5-8 minutes of editing per note — are the clinic's evidence, not a vendor claim. At minimum the restart either reproduces those savings with measurement or dies on evidence inside 30 days, either of which is progress over an unmeasured lapse.

Effort: Cost is inside budget: ~$198/month at two clinicians (pilot rate $99/clinician/month). A BAA was signed for the pilot; confirm it still holds and add scribe consent wording to intake. Verify the WebPT integration question first (owner's own condition). Adoption risk is Mark, mitigated structurally by making him the note auditor; Jessica's involvement is minimal by design.

Validation: Owner's standard: at least 30 minutes a day saved, edits under 10 minutes per note, notes signed same-day, no uptick in patient complaints about the room, five notes spot-

> audited weekly by Mark. If the WebPT integration answer is negative, proceed copy-paste only if the math still clears.

Referral intelligence: clean source data, monthly referrer report, physician feedback loop

Problem: Physician referrals are ~45-50 percent of new patients, but the WebPT source field is free text typed five different ways, so referrer performance cannot be reported; a sports-med referrer dropping from 6-8 patients a month to 1-2 went unnoticed for ~3 months; and communication with referrers is one-directional (ad hoc faxed discharge summaries, no systematic progress updates). Hendersonville, the expansion site, has only 5-6 active referrers and is the demand side of the expansion thesis.

Why it matters: You cannot grow — or protect — a referral base you cannot see. The owner's own words: 'I can't grow a referral base I can't see.' Catching a quiet referrer in weeks instead of months protects roughly half of new-patient volume, and a monthly report plus progress notes gives the owner a concrete reason for the referrer visits his freed hours are meant to fund.

AI approach: AI-assisted one-time normalization of the historical free-text source entries into a structured referrer list, a structured source picker at intake going forward, and AI-drafted quarterly progress-note templates that staff personalize and send to the top referrers — the reporting itself is a WebPT report run on a schedule.

How: Front desk selects the referrer from a dropdown instead of typing; a monthly referrer scorecard (volume by referrer, 6-month trend) flags anyone down materially; the owner acts on flags within two weeks (call or coffee); the top ten referrers get a short quarterly progress letter. Verify whether WebPT custom fields/reports support this before building anything outside it.

Impact: Mechanism: the Patel case cost an estimated 4-6 patients a month for a quarter before anyone noticed. Early detection plus top-ten feedback is the cheapest defense of a channel that supplies ~half of new patients, and it directly supports the Hendersonville referral development the expansion depends on.

Effort: Mostly process, not software: dropdown hygiene, a calendar-held monthly reporting ritual, and letter templates. The AI-assisted historical cleanup is a one-time job; verify WebPT custom-field and reporting capabilities first. Owner's time is the real cost — the visits and calls after each monthly report.

Validation: Percent of new patients with a structured (non-free-text) source after 30 and 60 days; the first monthly report produced within 30 days; every flagged quiet referrer contacted within two weeks; referral volume trend for the contacted practices over the following quarter.

Automated eligibility, benefits, and authorization pre-checks before the first visit

Problem: Verification takes 20-30 minutes per new patient in batched payer-portal sessions (about 90 new patients a month, so roughly 30-45 staff hours); three wrong-benefit quotes occurred last quarter (payers returned gym/wellness benefits instead of outpatient PT); and 2-3 workers-comp authorizations arrive late each month. Errors are found at or after the first visit, producing rework, patient confusion, and denials — eligibility/authorization is a top denial category at the outsourced biller (6-8 percent first-pass denials, ~34 days AR).

Why it matters: Fixes the problem upstream where it starts: correct copays quoted before the visit reduce the insurance-confusion losses among inquiries, fewer late authorizations reduce denied visits, and cutting verification toward the owner's sub-10-minute standard frees Marta and front-desk hours on both sites.

AI approach: Automated eligibility/benefits checks run ahead of scheduling with AI-assisted review of flagged exceptions (wrong-service benefits, auth-required plans), plus a workers-comp authorization status tracker with due-date alerts — humans handle exceptions, not the whole queue.

How: New-patient booking triggers an automated benefits check; results land in a queue Marta reviews; mismatches or auth-required plans are resolved before the first visit. The owner correctly framed the gating unknown: whether this needs a clearinghouse hook into WebPT billing is exactly the integration question to answer before believing it.

Impact: Mechanism: ~90 new patients a month times 20-30 minutes is the addressable time, and the measured error rate (3 wrong quotes last quarter, 2-3 late WC auths a month) is the addressable rework. Presently unquantified in dollars — the interview did not establish the revenue lost to resulting denials.

Effort: The most integration-dependent of the five: requires verifying a clearinghouse/WebPT billing integration path before commitment, which the interview did not investigate. Changes Marta's workflow from batch portal sessions to exception review.

Validation: Verification minutes per new patient (target <10), wrong-benefit quote incidents (target zero), late WC authorizations per month (from 2-3 toward zero), and first-pass denial rate (from 6-8 percent). Keep the outsourced billing arrangement unchanged.

Systematic discharge-to-reactivation engine: review asks, reactivation sends, Fit-for-Life relaunch

Problem: Only about a third of discharged patients are asked for a review despite strong ratings (148 at 4.8 Asheville; 41 at 4.7 Hendersonville); past patients and word of mouth supply ~20-25 percent of new patients but reactivation is unowned — a single 2024 email to ~300 past patients produced 4 bookings and was never repeated, ~1,800 past-patient emails sit unused in WebPT, and the Fit-for-Life membership has 11 members with no owner and no promotion.

Why it matters: Discharged patients are the cheapest demand source the clinic already owns, and the single small test (4 bookings from ~300 emails — pilot-scale evidence, not proof) plus 11 unpromoted members both suggest latent response. Reactivation sends can also target midday-friendly and pelvic-wellness audiences, feeding two capacity problems at once.

AI approach: AI-drafted, staff-reviewed message sequences: a discharge-triggered review ask and Fit-for-Life pitch added to the discharge workflow, a monthly reactivation send segmented from the WebPT email list, and templated check-ins at 90 days for appropriate discharged patients.

How: Add the ask to the discharge checklist so it happens every time; Jessica owns a monthly send calendar; measure per-send response. Consent and HIPAA-safe messaging practices apply (marketing vs treatment communications differ — confirm consent status of the existing list before first send, which the interview did not establish).

Impact: Mechanism: more consistent asks should raise review velocity (protecting the Google channel that supplies ~18 percent of new patients) and the list is large enough that even the 2024 response rate, repeated monthly, adds bookings. Magnitude is genuinely unknown — one small prior test only.

Effort: Low cost, mostly Jessica's calendar time plus AI-drafted copy; verify WebPT's email tooling and the consent status of the ~1,800-address list before the first send. Depends on therapists and front desk actually adding the ask to discharge.

Validation: Percent of discharges receiving the ask (target near 100 percent), reviews added per month, bookings per reactivation send, and Fit-for-Life membership (from 11). Give it one full quarter before judging.

Prioritized Top 3

#1 — Automated waitlist and cancellation filling with two-way texting

Why first: The owner reordered the draft ranking himself: texting goes first because it is revenue now ('a filled cancellation is cash in the door this week' — value), Jessica can own it inside her 4-5 weekly hours (feasibility), and it asks no clinician to change behavior — trading off against the scribe's proven-but-unverified-integration savings and the referral work's slower burn (time to benefit).

First step: Verify the WebPT calendar integration with the shortlisted tools using a real test calendar, write the intake consent script, and capture a two-week baseline of fill rate and no-show rate before the first text goes out.

Value: High: the owner named fill rate the fastest P&L lever; ~35 waitlisted patients, roughly half of prime-time cancels unfilled same-day, 11-12 percent leakage at ~$115 per visit.

Feasibility: Medium: no clinician behavior change and Jessica can run it, but WebPT calendar integration is unverified and the consent script must precede the first text.

Cost: Low: inside the stated $500-800/month tool budget; no one-time buildout.

Risk: Medium: the owner's own kill criteria are opt-out spikes and calendar sync that fights WebPT; both are visible in the first month.

Time to benefit: Weeks: fill rate and no-show movement are measurable within the first month of texting.

#2 — Measured AI scribe restart with a named note auditor

Why first: Second because the savings are already proven in the clinic's own pilot (value: 35-45 minutes/day at $198/month for two clinicians — cost: Low), but it requires clinician adoption (Mark) and an unverified WebPT integration (feasibility/risk), so it follows the texting rollout and starts in January once Jessica's rollout bandwidth frees up.

First step: Get the WebPT integration/API answer in writing, document a one-week baseline of documentation minutes per clinician, and re-confirm the BAA before relaunching with two clinicians and Mark as weekly note auditor.

Value: High: 35-45 minutes per clinician per day measured in the clinic's own six-week pilot across nine clinicians, plus the owner's freed hours fund gates-critical work (referrers, hiring, expansion prep).

Feasibility: Medium: savings proven locally and a BAA already signed, but the WebPT integration question was never asked and the first pilot's copy-paste workflow annoyed clinicians.

Cost: Low: ~$198/month at two clinicians at the pilot's $99/clinician/month rate; verification effort before relaunch is staff time only.

Risk: Medium: the first pilot lapsed without an owner or measurements; adoption risk concentrated in Mark is mitigated by making him the note auditor.

Time to benefit: Weeks to months: a 30-day measured restart with the owner's keep-it standard (>=30 min/day, edits <10 min/note, same-day signatures).

#3 — Referral intelligence: clean source data, monthly referrer report, physician feedback loop

Why first: Third on speed of benefit but not on importance: feasibility is High and cost is Low (mostly process hygiene and a scheduled report), and the owner corrected the draft to weight it toward Hendersonville, where the thinner referral base (5-6 active referrers) is the demand-side risk to the entire expansion thesis — the tradeoff is that referral relationships pay back over months, not weeks.

First step: Run the one-time AI-assisted cleanup of the free-text source field, switch intake to a structured referrer picker, and produce the first monthly referrer report within 30 days.

Value: Medium: defends a channel supplying ~45-50 percent of new patients (the quiet-referrer case cost an estimated 4-6 patients a month for a quarter) and underpins Hendersonville demand development; harder to tie to immediate dollars.

Feasibility: High: dropdown hygiene, a monthly scheduled WebPT report, and letter templates — verify WebPT custom-field/report support before building outside it.

Cost: Low: one-time AI-assisted cleanup plus staff process; the owner's referrer visits are the main ongoing cost.

Risk: Low: worst case is process discipline slips; no patient-facing or clinical risk. Data quality improves monotonically once the picker is in place.

Time to benefit: Months: quiet-referrer detection improves within the first two monthly reports; referrer relationship effects compound over a quarter.

30-Day Action Plan

A 30-day starting plan that respects the owner's two hard constraints — Jessica's 4-5 weekly hours and one rollout at a time. Week 1: capture baselines (two-week fill rate and no-show rate, one-week documentation minutes per clinician, verification minutes per new patient); ask WebPT two direct questions — texting/calendar integration options for scheduling tools and scribe/API integration options; draft the intake consent script for texting. Week 2: run tool demos against a real test calendar and choose the texting tool; switch intake to the structured referrer picker and start the one-time AI-assisted cleanup of the free-text source field; begin the texting pilot at both sites with Jessica watching opt-outs daily. Week 3: produce the first monthly referrer report and contact any flagged quiet referrer within two weeks; shortlist referrer visits for the top ten practices (Hendersonville primary care first); define the scribe restart protocol — two

clinicians, Mark as auditor of five notes weekly, the owner's keep-it metrics on one page. Week 4: first texting checkpoint against the 90-day standard (>=60 percent of open slots filled within two business days; no-shows -3 points; opt-outs acceptable) — keep, adjust, or kill; send the first quarterly progress notes to top referrers; schedule the scribe restart for January so the two rollouts never land on Jessica in the same month. In parallel, background hygiene the owner named himself: reconcile the 11-12 percent no-show figure, the ~$105 vs ~$120 per-visit gap between sites against payer and contract mix, and put basic attribution on the ~$1,400/month (approximate) Google Ads spend. None of this changes the expansion gates: the lease decision still hinges on a signed hire and the $85,000/month run rate, and the cash-reserve gap (~$95,000 held vs ~$120,000 gate) deserves the same monthly attention the metrics list gets.

Risks & Metrics

- WebPT calendar integration for waitlist/texting tools is not established — the first pilot's copy-paste history shows why this must be verified before contracting.
- WebPT scribe integration/API path is not established; the clinic never asked, and the 2026 pilot ran on manual copy-paste.
- Not investigated in this interview: eligibility/benefits clearinghouse integration options for WebPT billing (gates opportunity 4 entirely).
- Not investigated in this interview: the competitor landscape in Asheville/Hendersonville and its effect on referral and pricing behavior.
- Not investigated in this interview: front-desk phone system capabilities (whether the current VoIP setup supports call tracking or texting) — any assumption here is unverified.
- Not investigated in this interview: the consent status of the ~1,800-address past-patient email list for marketing reactivation sends.
- The no-show rate is quoted as a range (11-12 percent) and the per-visit gap between sites (~$105 vs ~$120) is not reconciled against payer and employer-contract mix — reconcile both before using either as a pilot baseline.
- Google Ads spend (~$1,400/month, owner-stated approximation) continues without attribution to booked evaluations; the 10-15 calls/month figure is an estimate.
- Clinician composition was corrected by the owner in draft review: seven physical therapists and two PTAs (the first draft miscounted all nine as physical therapists).

- Reactivation and Fit-for-Life economics rest on one small test (4 bookings from ~300 emails) and 11 unpromoted members — pilot-scale evidence, not proof.
- Fill rate: percent of open slots filled within two business days (target >=60 percent; owner's keep-it standard).
- No-show + late-cancel rate (baseline 11-12 percent, target -3 points).
- Texting opt-out rate (owner's kill criterion — watch from day one).
- Documentation minutes per clinician per day (baseline 45-60; target <30).
- Note edit time per note (<10 minutes) and same-day signature rate during the scribe restart.
- Verification minutes per new patient (baseline 20-30; target <10) and wrong-benefit quote incidents (3 last quarter).
- Late workers-comp authorizations per month (baseline 2-3).
- First-pass denial rate (baseline 6-8 percent) and AR days (~34).
- Percent of new patients with a structured referral source; flagged quiet referrers contacted within two weeks.
- Reviews added per month, bookings per reactivation send, and Fit-for-Life membership (baseline 11).